

Práctica 4: Orquestación de Contenedores en AWS ECS

Introducción

En esta práctica, se ha implementado un servicio en **AWS Elastic Container Service (ECS)** utilizando **AWS Fargate** como infraestructura sin servidor (*serverless*). Se documenta el proceso de creación del clúster, la definición de tareas, el balanceador de carga y la creación del servicio.

Además, se ha experimentado un error durante la creación del clúster debido a un problema interno de AWS (**Error Code: ServerException, Status Code: 500**). La solución ha sido evaluar en el stack el problema y relanzar manualmente el proceso.

1. Creación del Clúster

Pasos realizados:

1. Acceder a **AWS ECS** y seleccionar "**Create Cluster**".
 2. Configurar:
 - **Nombre:** `alba-digi-cluster`
 - **Infraestructura:** `AWS Fargate (serverless)`
 - **Tags:** Key: `Name`, Value: `alba-digi-cluster`
 3. Crear el clúster.
 4. Debido al error **500 ServerException**, AWS CloudFormation reintentó la creación del clúster.
 5. Una vez completada la creación, se verificó el estado en **CloudFormation > Events**, confirmando el evento **CREATE_COMPLETE** para `ECSCluster`.
 6. Se capturó la pantalla del evento exitoso como evidencia (adjunta a este documento).
-

2. Creación de la Definición de Tarea (Task Definition)

1. Acceder a **AWS ECS** y seleccionar "**Task Definitions**".
2. Crear una nueva definición de tarea y seleccionar "**Fargate**".
3. Configurar:
 - **Nombre:** `alba-digi-task`
 - **Task Role:** `LabRole`
 - **Network Mode:** `awsvpc`
 - **Sistema Operativo:** `Linux`
 - **Task Execution Role:** `LabRole`

- **CPU:** 1 vCPU
 - **Memoria:** 3GB
 - 4. Agregar un contenedor:
 - **Nombre:** alba-digi-api-rest
 - **Imagen:** public.ecr.aws/n0h7v2z5/alba-digi-api-rest:latest
 - **Memoria (Soft Limit):** 2GB
 - **Puerto:** 5050 (TCP, HTTP)
 - 5. Crear la definición de tarea y capturar una captura de pantalla.
-

3. Creación del Balanceador de Carga

1. Acceder a **AWS EC2** y seleccionar "**Load Balancers**".
 2. Crear un nuevo balanceador de carga **Application Load Balancer**.
 3. Configurar:
 - **Nombre:** alba-digi-load-balancer
 - **Scheme:** Internet-facing
 - **IP Address Type:** IPv4
 - **Network Mapping:** Todas las zonas disponibles
 - **Security Groups:** Dejar el predeterminado
 4. Crear un **Target Group**:
 - **Nombre:** alba-digi-target-group
 - **Protocolo:** HTTP, **Puerto:** 5050
 - **Health Check:** HTTP, Path: /alive
 5. Asociar el *Target Group* al balanceador y completar la configuración.
 6. Capturar la pantalla del balanceador de carga creado.
-

4. Creación del Servicio

1. Acceder a **AWS ECS** y seleccionar el clúster alba-digi-cluster.
2. Crear un servicio y configurar:
 - **Launch Type:** Fargate
 - **Task Definition:** alba-digi-task (latest)
 - **Nombre del Servicio:** alba-digi-service
 - **Number of tasks:** 2
 - **Load Balancer Type:** Application Load Balancer
 - **Health Check Grace Period:** 30
 - **Container to Load Balance:** alba-digi-api-rest:5050:5050
3. Completar la configuración y capturar una imagen del servicio creado.

5. Prueba del Servicio

1. Copiar el **DNS Name** del balanceador de carga desde **EC2 > Load Balancers**.

En una terminal, ejecutar:

```
curl -v http://{DNS_NAME}/alive
```

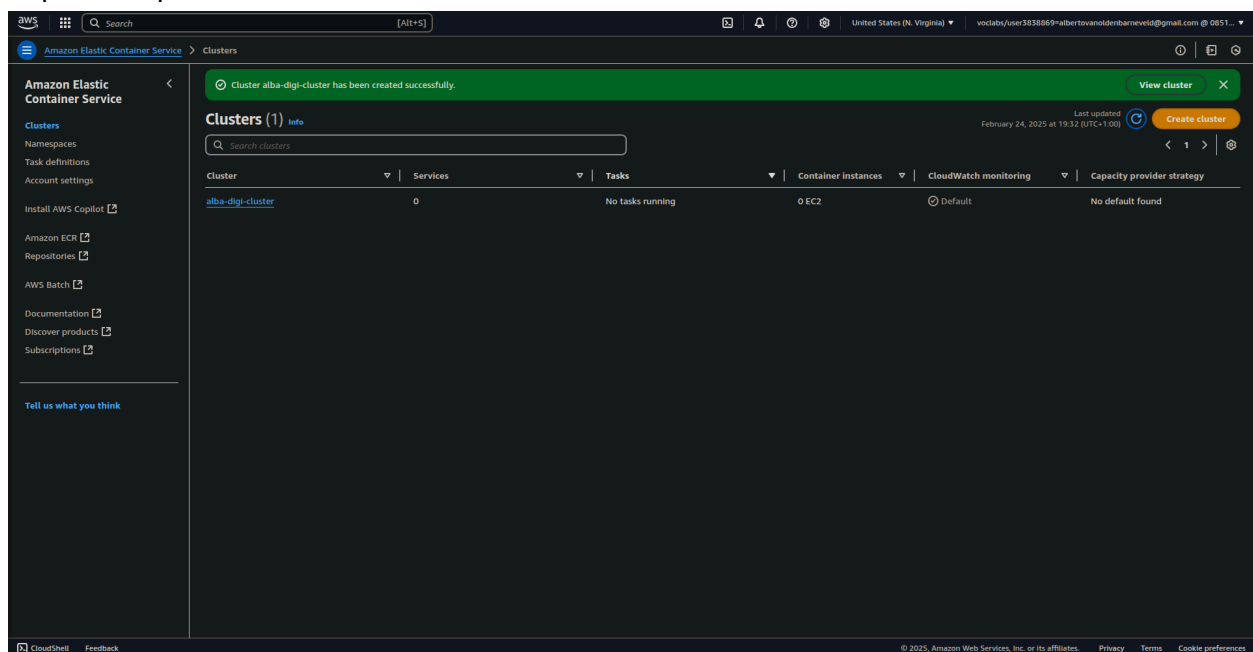
2. (Reemplazar **{DNS_NAME}** con la URL copiada).
 3. Si la petición falla, ajustar las **Inbound Rules** del Security Group para permitir conexiones HTTP.
 4. Volver a ejecutar **curl** y capturar la pantalla de la respuesta exitosa.
-

Conclusión

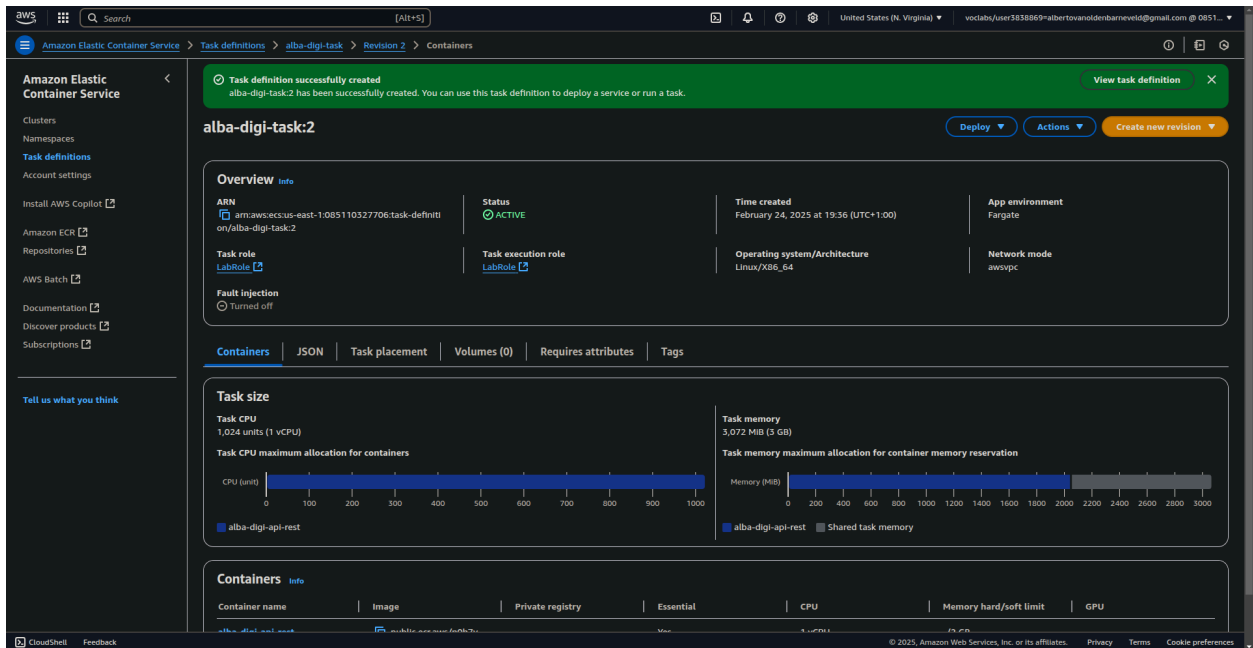
Se ha desplegado un servicio en **AWS ECS con Fargate**, configurando un **clúster**, una **definición de tarea**, un **balanceador de carga**, y un **servicio** con dos tareas. A pesar del error inicial en la creación del clúster, la solución fue esperar a que AWS CloudFormation reintentara el proceso.

Las capturas de pantalla tomadas a lo largo del proceso validan la correcta ejecución de cada etapa.

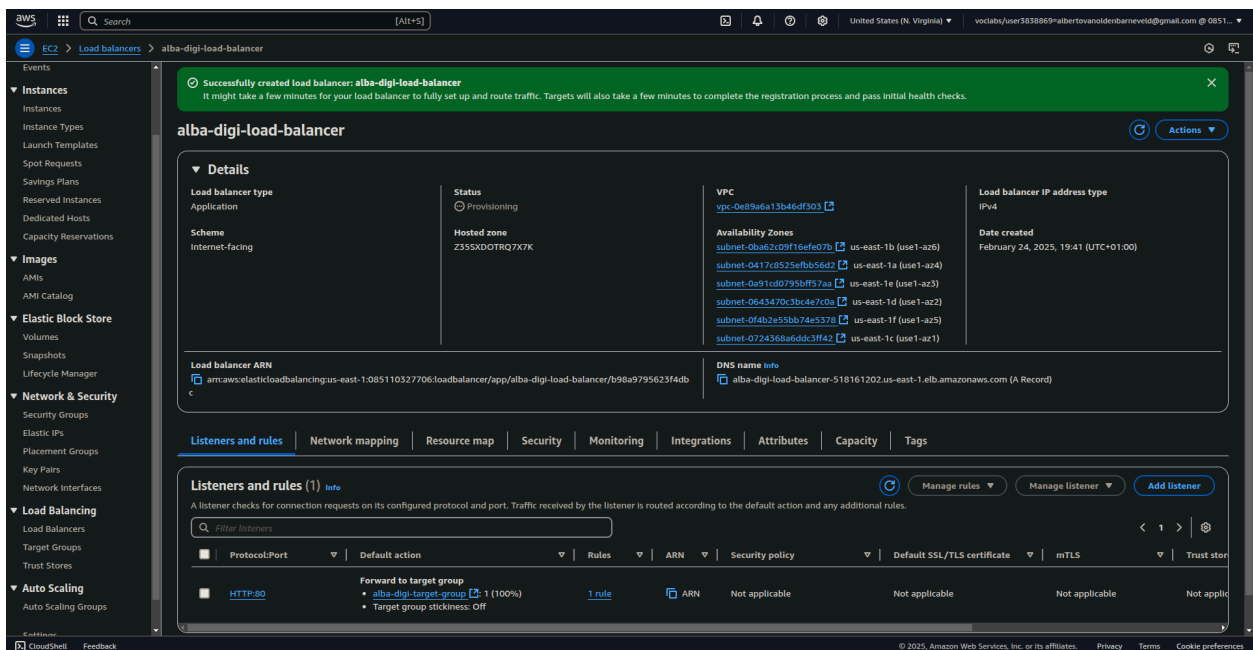
Captura de pantalla 1:



Captura de pantalla 2:



Captura de pantalla 3:



Capturas de pantalla 4:

Create [info](#)

Environment AWS Fargate

Existing cluster

alba-digi-cluster

▼ Compute configuration (advanced)

Compute options [info](#)

To ensure task distribution across your compute types, use appropriate compute options.

☐ Capacity provider strategy

Specify a launch strategy to distribute your tasks across one or more capacity providers.

☒ Launch type

Launch tasks directly without the use of a capacity provider strategy.

Launch type [info](#)

Select either managed capacity (Fargate), or custom capacity (EC2 or user-managed, External instances). External instances are registered to your cluster using the ECS Anywhere capability.

FARGATE

Platform version [info](#)

Specify the platform version on which to run your service.

LATEST

Deployment configuration

Application type [info](#)

Specify what type of application you want to run.

☒ Service

Launch a group of tasks handling a long-running computing work that can be stopped and restarted. For example, a web application.

☐ Task

Launch a standalone task that runs and terminates. For example, a batch job.

Task definition

Select an existing task definition. To create a new task definition, go to [Task definitions](#).

☒ Specify the revision manually

Manually input the revision instead of choosing from the 100 most recent revisions for the selected task definition family.

Family

alba-digi-task



Revision

2 (LATEST)

Service name

Assign a service name that is unique for this cluster.

alba-digi-service

Up to 255 letters (uppercase and lowercase), numbers, underscores, and hyphens are allowed. Service names must be unique within a cluster.

Service type [info](#)

Specify the service type that the service scheduler will follow.

☒ Replica

Place and maintain a desired number of tasks across your cluster.

☐ Daemon

Place and maintain one copy of your task on each container instance.

Desired tasks

Specify the number of tasks to launch.

2

Availability Zone rebalancing [info](#)

☒ Turn on Availability Zone rebalancing

Amazon ECS automatically detects Availability Zone imbalances in task distributions across an ECS service, and evenly redistributes ECS service tasks across Availability Zones.

▼ Deployment options

Deployment type | [Info](#)

Select a deployment controller type for the service.

☒ Rolling update☐ Blue/green deployment (powered by AWS CodeDeploy)**Min running tasks %** | [Info](#)

Specify the minimum percent of running tasks allowed during a service deployment.

values 0 to %

Max running tasks % | [Info](#)

Specify the maximum percent of running tasks allowed during a service deployment.

values 0 to %

▼ Deployment failure detection | [Info](#)

⚠ Deployment protection settings have been turned off.

☐ Use the Amazon ECS deployment circuit breaker

If the service can't reach a steady state because a task failed to launch, the deployment fails.

☐ Use CloudWatch alarm(s)

If the CloudWatch alarm or alarms that you specify transition to the ALARM state, the deployment fails.

▼ Service Connect - optional | [Info](#)

Service Connect allows for service-to-service communications with automatic discovery using short names and standard ports.

☐ Use Service Connect

Configure the namespaces, and the services to interconnect.

► Service discovery - optional

Service discovery uses Amazon Route 53 to create a namespace for your service, which allows it to be discoverable via DNS.

► Networking**▼ Load balancing - optional**

Configure load balancing using Amazon Elastic Load Balancing to distribute traffic evenly across the healthy tasks in your service.

☒ Use load balancing**Load balancer type** | [Info](#)

Specify the load balancer type to distribute incoming traffic across the tasks running in your service.

☒ Application Load Balancer

An Application Load Balancer makes routing decisions at the application layer (HTTP/HTTPS), supports path-based routing, and can route requests to one or more ports.

☐ Network Load Balancer

A Network Load Balancer makes routing decisions at the transport layer (TCP/UDP).

Container

The container and port to load balance the incoming traffic to



Host port:Container port

Application Load Balancer
Specify whether to create a new load balancer or choose an existing one.

☐ Create a new load balancer
☒ Use an existing load balancer

Load balancer
Select the load balancer you wish to use to distribute incoming traffic across the tasks running in your service.

alba-digi-load-balancer

Health check grace period [info](#)
30 seconds

Listener [info](#)
Specify the port and protocol that the load balancer will listen for connection requests on.

☐ Create new listener
☒ Use an existing listener

Listener 80:HTTP

Listener rules for 80:HTTP [\(1\)](#)
Traffic received by the listener is routed according to its rules. Rules are evaluated in priority order, from the lowest value to the highest value. The default rule is evaluated last.

Evaluation order	Rule path	Target group
default	/	alba-digi-target-group ↗

Target group [info](#)
Specify whether to create a new target group or choose an existing one that the load balancer will use to route requests to the tasks in your service.

☐ Create new target group
☒ Use an existing target group

Target group name alba-digi-target-group

Health check path /alive

Health check protocol HTTP

► **VPC Lattice - optional** [info](#)
Fully managed application networking service to connect, secure, and monitor your services across multiple accounts and virtual private clouds (VPCs). When you use VPC Lattice, there is a cost associated with it.

► **Service auto scaling - optional**
Automatically adjust your service's desired count up and down within a specified range in response to CloudWatch alarms. You can modify your service auto scaling configuration at any time to meet the needs of your application.

► **Volume - optional** [info](#)
Configure a data volume to provide additional storage for the containers in the task.

Capturas de pantalla 5: Tareas del servicio

alba-digi-service [info](#)

Status Active | Tasks (2 Desired) 0 Pending | 2 Running | Task definition: revision [alba-digi-task:2](#) | Deployment status Success

Health and metrics | **Tasks** | Logs | Deployments | Events | Configuration and networking | Service auto scaling | Tags

Tasks (1/2)

Filter desired status: Any desired status | Filter launch type: Any launch type

Task	Last status	Desired sta...	Task definit...	Health status	Started by	Started at	Container instan...	Launch type	Platform versi...	CPU	Memory
4b5bd743...	Running	Running	alba-digi-task:2	Unknown	ecs-svc/49630608980...	1 minute ago	-	FARGATE	1.4.0	1 vCPU	3 GB
519a778f...	Running	Running	alba-digi-task:2	Unknown	ecs-svc/49630608980...	1 minute ago	-	FARGATE	1.4.0	1 vCPU	3 GB

Containers for task 4b5bd74373514c518ccc94f1b788e15e

Containers (1)

Container name	Container runtime ID	Image URI	Image Digest	Status	Health status	CPU	Memory hard/soft limit
alba-digi-api-test	4b5bd74373514c...	public.ecr...	sha256:f9...	Running	Unknown	1 vCPU	- / 2 GB

Eventos del servicio:

aws [Search] [Alt+S] United States (N. Virginia) voclabs/user3838069=albertovandenbarneveld@gmail.com @ 0851...

Amazon Elastic Container Service > Clusters > alba-digi-cluster > Services > alba-digi-service > Events

alba-digi-service info

Last updated February 24, 2025 at 19:52 (UTC+1:00) [Update service](#) [Delete service](#)

Service overview info

Status
 Active

Tasks (2 Desired)
 0 Pending | 2 Running

Task definition: revision
 alba-digi-task-2

Deployment status
 Success

Health and metrics | Tasks | Logs | Deployments | **Events** | Configuration and networking | Service auto scaling | Tags

Events (5) info

Started at	Message	Event ID
February 24, 2025 at 19:51 (UTC+1:00)	service alba-digi-service has reached a steady state.	4a11d8ff-f341-4f7a-b101-cae24fa1512d
February 24, 2025 at 19:51 (UTC+1:00)	service alba-digi-service deployment ecs-svc/4963060898094496537 deployment completed.	40a87543-f0f7-42b3-a467-d34297ede251
February 24, 2025 at 19:50 (UTC+1:00)	service alba-digi-service registered 1 targets in target-group alba-digi-target-group 🔗	832be845-5a7a-4b30-aa57-4b2feac5d164
February 24, 2025 at 19:50 (UTC+1:00)	service alba-digi-service has started 1 tasks: task b19a778f352545398e91af23cf239fc0 .	f6be3415-9883-44d5-8fe7-48674894abbf
February 24, 2025 at 19:50 (UTC+1:00)	service alba-digi-service has started 1 tasks: task 4b5bd74373514c518ccc94f1b780e15a .	9ec4f20f-488b-445d-93b2-e1155f03deeb

<https://us-east-1.console.aws.amazon.com/ecs/v2/clusters/alba-digi-cluster/services/alba-digi-service/events?region=us-east-1>

© 2025, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

```
Mon Feb 24 19:59
alberto@alberto-IdeaPad-3-15ITL6: ~/Desktop

alberto@alberto-IdeaPad-3-15ITL6:~/Desktop$ curl -v http://alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com/alive
* Host alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com:80 was resolved.
* IPv6: (none)
* IPv4: 52.5.120.238, 34.194.89.57, 52.22.26.91
*   Trying 52.5.120.238:80...
* connect to 52.5.120.238 port 80 from 192.168.0.65 port 45514 failed: Connection timed out
*   Trying 34.194.89.57:80...
* ipv4 connect timeout after 82781ms, move on!
*   Trying 52.22.26.91:80...
^C
alberto@alberto-IdeaPad-3-15ITL6:~/Desktop$ curl -v http://alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com/alive
* Host alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com:80 was resolved.
* IPv6: (none)
* IPv4: 52.5.120.238, 34.194.89.57, 52.22.26.91
*   Trying 52.5.120.238:80...
* Connected to alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com (52.5.120.238) port 80
> GET /alive HTTP/1.1
> Host: alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com
> User-Agent: curl/8.5.0
> Accept: */*
>
< HTTP/1.1 200 OK
< Date: Mon, 24 Feb 2025 18:59:13 GMT
< Content-Type: text/plain; charset=UTF-8
< Content-Length: 63
< Connection: keep-alive
<
* Connection #0 to host alba-digi-load-balancer-518161202.us-east-1.elb.amazonaws.com left intact
/ - Hello alive! Host:ip-172-31-93-57.ec2.internal/172.31.93.57alberto@alberto-IdeaPad-3-15ITL6:~/Desktop$
```